

TALLINN UNIVERSITY OF TECHNOLOGY
School of Information Technologies
Thomas Johann Seebeck Department of Electronics

Gaspar Karm, IVEM153410

Pyha

Master's Thesis

Supervisors:

Muhammad Mahtab Alam
PhD

Yannick Le Moullec
PhD

Tallinn 2017

TALLINNA TEHNIKAÜLIKOOL
Infotehnoloogia teaduskond
Thomas Johann Seebecki elektroonikainstituut

Gaspar Karm, IVEM153410

Pyha

Magistritöö

Juhendajad:

Muhammad Mahtab Alam
PhD

Yannick Le Moullec
PhD

Tallinn 2017

Author's declaration of originality

I hereby certify that I am the sole author of this thesis. All the used materials, references to the literature and the work of others have been referred to. This thesis has not been presented for examination anywhere else.

Author: Gaspar Karm
[pp.kk.aaaa]

Abstract

The thesis is in English and contains [pages] pages of text, 5 chapters, [figures] figures, [tables] tables.

Annotatsioon

Lõputöö on kirjutatud inglise keeles ning sisaldab teksti [lehekülgede arv töö põhiosas] leheküljel, 5 peatükki, [jooniste arv] joonist, [tabelite arv] tabelit.

Contents:

| | | |
|----------|--|-----------|
| 1 | Introduction | 1 |
| 1.1 | Problem statement | 2 |
| 1.2 | Structure | 3 |
| 2 | Hardware design with Pyha | 5 |
| 2.1 | Introduction | 5 |
| 2.2 | Sequential logic | 8 |
| 2.2.1 | Accumulator | 8 |
| 2.2.2 | Block processing and sliding adder | 10 |
| 2.3 | Fixed-point designs | 13 |
| 2.3.1 | Basics | 13 |
| 2.3.2 | Fixed-point sliding adder | 14 |
| 2.4 | Summary | 15 |
| 3 | Conversion to VHDL | 17 |
| 3.1 | Sequential, Object-oriented style for VHDL | 18 |
| 3.1.1 | Defining registers with variables | 19 |
| 3.1.2 | Creating instances | 21 |
| 3.1.3 | Final OOP model | 22 |
| 3.1.4 | Examples | 23 |
| 3.1.5 | Multiple clock domains | 24 |
| 3.2 | Converting Python to VHDL | 25 |
| 3.2.1 | Finding the types | 26 |
| 3.2.2 | Syntax conversion | 27 |
| 3.3 | Summary | 28 |
| 4 | Case studies | 29 |
| 4.1 | Moving average filter | 29 |
| 4.2 | Linear-phase DC removal Filter | 32 |
| 4.3 | Comparison to other tools | 34 |
| 5 | Summary | 37 |
| | Bibliography | 39 |

Chapter 1

Introduction

Main tools today for designing digital hardware are VHDL and SystemVerilog (SV). SV is aggressively promoted by the big EDA (Cadence, Mentor, Synopsys) along with Universal Verification Methodology (UVM). At 2003, Aart de Geus, Synopsys CEO, has stated that SV will replace VHDL in 10 years [1]. It is true that tool vendors have stopped enchanting VHDL support, even so the development is going on in the open-source sphere, where VHDL-2017 standard [2] is in the development. In addition, active development is going on open source simulator GHDL, Open Source VHDL Verification Methodology (OSVVM) [3] and unit-testing library VUnit [4]. All the improvements the traditional languages receive aim to ease the verification task, the synthesizable parts of VHDL and SV have stayed mostly the same for the past 10 years.

Numerous projects exist that propose to use higher level HDL languages in order to raise the abstraction level. For example, MyHDL turns Python into a hardware description and verification language [5]. Or CλaSH [6], purely Haskell based functional language developed at University of Twente. Recently Chisel [7] has been gaining some traction, it is an hardware construction language developed at UC Berkeley that uses Scala programming language, providing functional and object-oriented features. Still none of these tools have seen widespread adaption.

On the other front, high-level synthesis(HLS) tools aim to automate the refinement from the algorithmic level to RTL [8]. Lately the Vivado HLS, developed by Xilinx, has been gaining popularity. As of 2015, it is included in the free design suite of Vivado (device limited). Problem with HLS tools is that they are often promoted as direct C to RTL tools but in reality often manual code transformations and guidelines are needed, in order to archive reasonable performance [9]. The designer must already know how the RTL works in order to give these instructions.

The DSP systems can be described in previously mentioned HLS or HDL languages, but the most productive way is to use MATLAB/Simulink/HDLConverter flow, which allows users to describe their designs in Simulink or MATLAB and using HDL convertible blocks provided by MATLAB or FPGA tool vendor [10].

1.1 Problem statement

There is no doubt that MATLAB based workflow offers an highly productive path from DSP models to hardware. However these tools can easily cost over tens of thousands euros, often FPGA vendor tools are required that add additional annual cost [10]. Using these tools is not suitable for reproducible research and is completely unusable for open-source designs. Thus the designers must turn to alternative design flows, for example [11] provides an hardware implementation of an ADS-B (automatic dependent surveillance – broadcast). First, they did the prototyping in the MATLAB environment, the working model was then translated to C for real-time testing and fixed-point modeling. Lastly the C model was manually converted to VHDL.

This thesis introduces Pyha, a new Python based HDL developed during the masters thesis program, with an goal to provide open-source alternative for the MATLAB based flows. Pyha raises the RTL design abstraction by enabling sequential and object-oriented style. DSP systems can be built by using the fixed-point type and semi-automatic conversion from floating point. In addition, this work makes an effort to simplify the testing process of hardware systems by providing better simulation interface for unit-testing.

Basis of Pyha is Python, a general purpose programming language that is especially well suited for rapid prototyping and modeling. Python has also found its place in scientific projects and academia by offering most of what is familiar from MATLAB, free of charge. Scientist are already shifting from MATLAB to Python in order to conduct research that is reproducible and accessible by everyone [12]. Fig. 1.1 shows the popularity comparison (based on Google searches) of Python, MATLAB and C. Python is far ahead and the only one with positive trend.




Fig. 1.1: PYPL(PopularitY of Programming Language) [13]. Python 15.1%, C 6.9%, MATLAB 2.7%

Furthermore, this work introduces the sequential OOP VHDL model, that is developed to allow simpler conversion from Python to VHDL. Side contribution

1.2 Structure

This thesis is divided into 3 chapters. In chapter 1, main concepts of Pyha are introduced. Following chapter shows First chapter of this thesis gives an overview of the developed tool Pyha and how it can be used for hardware design. Follows the examples that show how Pyha can be used to relatively easily construct moving-average filter and by reusing it the DC-removal filter. Final chapter describes the one of the contributions of this thesis, the sequential VHDL OOP model and how Python is converted to it.

Chapter 2

Hardware design with Pyha

This chapter introduces the main contribution of this thesis, Pyha - a tool to design digital hardware in Python.

Pyha proposes to program hardware in the same way as software; much of this chapter is focused on showing differences between hardware and software constructs.

The first half of the chapter demonstrates how basic hardware constructs can be defined, using Pyha.

The second half introduces the fixed-point type and provides use-cases on designing with Pyha.

All the examples presented in this chapter can be found online [HERE](#), including all the Python sources, unit-tests, VHDL conversion files and Quartus project for synthesis.

Todo

organise examples to web and put link

2.1 Introduction

While the conventional HDL languages promote concurrent and entity oriented model, this is more confusing. In this thesis, Pyha has been designed as an sequential object-oriented language, that works directly on Python code. Using an sequential design flow is much easier to understand and is equally well synthesizable as shown by this thesis. Object-oriented design helps to better abstract the RTL details and ease design reuse.

For illustration purposes, [Listing 2.1](#) shows an example Pyha design. The `main` function has been chosen as a top level entry point, other functions can be used as pleased.

Note: First few examples of this chapter use `integer` types in order to reduce complexity.

Listing 2.1: Simple combinatory design, implemented in Pyha

```
class Basic(HW):
    def main(self, x):
        a = x + 1 + 3
        b = a * 314

        if a == 9:
            b = 0

        return a, b
```

One of the contributions of this thesis is sequential OOP VHDL model, that is used by Pyha for simple conversion to VHDL. Example of the VHDL conversion is shown on Listing 2.2, more details are given in chapter Section 3.

Listing 2.2: Listing 2.1 converted to VHDL, by Pyha

```
procedure main(self:inout self_t; x: integer; ret_0:out integer; ret_1:out_
↪integer) is
    variable a: integer;
    variable b: integer;
begin
    a := x + 1 + 3;
    b := a * 314;

    if a = 9 then
        b := 0;
    end if;

    ret_0 := a;
    ret_1 := b;
end procedure;
```

Fig. 2.1 shows the synthesis result. The `a` output is formed by adding ‘1’ and ‘3’ to the `x` input. Next the `a` signal is compared to 9, if equal `b` is outputted as 0, otherwise `b = a * 314`. That exactly complies with the Python and VHDL descriptions.

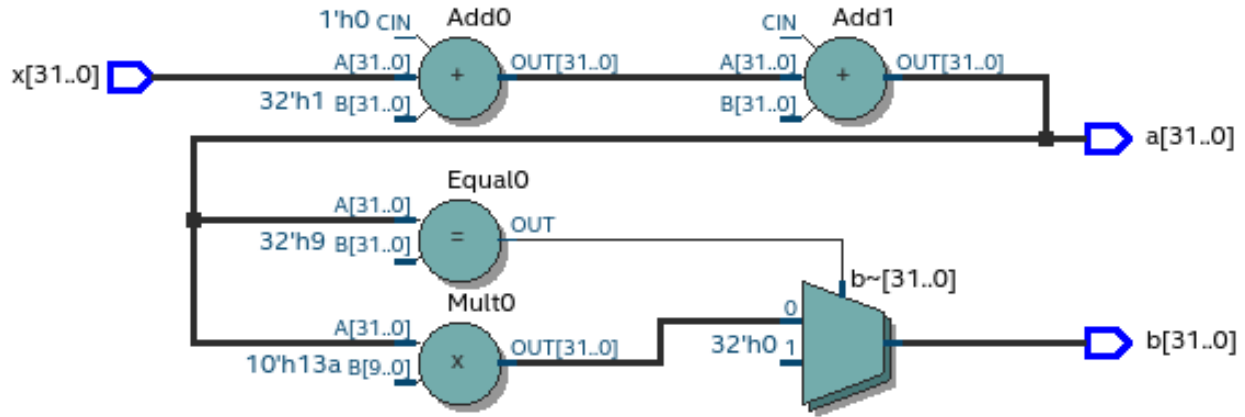


Fig. 2.1: Synthesised RTL of Listing 2.2 (Intel Quartus RTL viewer)

One aspect of hardware design that Pyha aims to improve is testing. Conventional tools like VHDL require the construction of special testbenches that can be executed on simulators. Even the higher level tools often don't simplify this step, for example the C based tools HLS tools want testbench in C language, which is not an improvement from VHDL or Verilog.

First of all, Pyha has been designed so that the synthesis output is behaviourally equivalent to the Python run output, this means that Pyha designs can use all the Python debugging tools. `add_multi_debug` shows a debugging session on the Listing 2.1 code, this can drastically help the development process.

```
def main(self, x): self: <main.Basic object at 0x7f2c21f7a630> x: 5
  a = x + 1 + 3 a: 9
  b = a * 314 b: 2826
  if a == 9:
    b = 0
  return a, b
```

Fig. 2.2: Debugging using PyCharm (Python editor)

Furthermore, unit testing is accelerated by providing `simulate(dut, x)` function, that runs the following simulations without any boilerplate code:

- Model: this can be any Python code that fits as an high level model;
- Pyha: like Listing 2.1, Python domain simulation;
- RTL: converts the Pyha model to VHDL and uses the combination of GHDL and Cocotb for simulation;
- GATE: synthesises the VHDL code, using Intel Quartus, and simulates the resulting gate-level netlist.

This kind of testing function enables test-driven development, where tests can be first defined for the model and fully reused for later RTL implementation. `pyha.adder_test` shows an

example unit test for the `Basic()` module.

Listing 2.3: Unit test for the Basic module

```
x = [1, 2, 3, 4, 5, 6, 7, 8]
dut = Basic()
y = simulation(dut, x)
# assert something
```

2.2 Sequential logic

In hardware registers are used as an memory element and for pipelining. In general digital logic synthesis relies on timing synthesis that only works when analized logic is between registers.

The way how registers are inferred is a fundamental difference between the RTL and HLS languages. Main complexity of HLS is about automatically inferring registers for memory elements or for pipelining. RTL languages on the other hand leave the task up to the designer. In this work, Pyha has been designed to follow the RTL language approach, because this comes free with conversion to VHDL. In future extensions can be considered.

In conventional programming, most commonly state is captured by using the class variables, which can keep the values between function calls. Inspired from this, all the class variables in Pyha are handled as registers, class functions can be interpreted as combinatory functions calculating the next state values for the registers.

2.2.1 Accumulator

Consider the design of an accumulator (Listing 2.4); it operates by sequentially adding up all the input values.

Listing 2.4: Accumulator implemented in Pyha

```
1 class Acc(HW):
2     def __init__(self):
3         self.acc = 0
4
5     def main(self, x):
6         self.next.acc = self.acc + x
7         return self.acc
```

The class structure in Pyha has been designed so that the `__init__` function shall define all the memory elements in the design, the function itself is not converted to VHDL, only the variables are extracted. For example `__init__` function could be used to call `scipy.signal`.

`firwin()` to design FIR filter coefficients, initial assignments to class variables are used as register initial/reset values.

Note the `self.next.acc = ...`, simulates the hardware behaviour of registers, that is delayed assignment. In general this is equivalent to the VHDL `<=` operator. Values are transferred from **next** to **current** just before the `main` call. In general Pyha abstracts the clock signal away by denoting that each call to `main` is a clock edge. Think that the `main` function is started with the **current** register values known and the objective of the `main` function is to find the **next** values for the registers.

The synthesis results shown in the Fig. 2.3 shows the adder and register.




Fig. 2.3: Synthesis result of Listing 2.4 (Intel Quartus RTL viewer)

One inconvenience is that every register on signal path delays the output signal by 1 sample, this is also called pipeline delay or latency. This situation is shown on Fig. 2.4 that shows the simulation results for the `Acc` module. Note that the model is implemented without register semantics, thus has no pipeline delays. This can be seen from the Fig. 2.4, hardware related simulations are delayed by 1 compared to the software model.




Fig. 2.4: Simulation of the `Acc` module, input is a random integer `[-5;5]`

Pyha reserves a `self._delay` variable, that hardware classes can use to specify their delay. Simulation functions read this variable and compensate the simulation data so that the delay is compensated, so that the compensation does not have to be made in unit-tests. Setting the `self._delay = 1`` in the `__init__` function would shift the hardware simulations left by 1 sample, so that all the simulations would be exactly equal.

2.2.2 Block processing and sliding adder

Common technique required to implement DSP systems is block processing i.e. calculating results on a block of input samples. Until now, the `main` function has worked with a single input sample, registers can be used to keep history of samples, so that block processing can be applied.

For an example, consider an algorithm that adds the last 4 input values (Listing 2.5). Listing 2.5 shows an implementation that keeps track of the last 4 input values and sums them. Note that the design also uses the output register `y`.

Listing 2.5: Sliding adder algorithm

```
class SlidingAdder(HW):
    def __init__(self):
        self.shr = [0, 0, 0, 0] # list of registers
        self.y = 0

    def main(self, x):
        # add new 'x' to list, throw away last element
        self.next.shr = [x] + self.shr[:-1]

        # add all element in the list
        sum = 0
        for a in self.shr:
            sum = sum + a

        self.next.y = sum
        return self.y
```

The `self.next.shr = [x] + self.shr[:-1]` line is also known as a ‘shift register’, because on every call it shifts the list contents to the right and adds new `x` as the first element. Sometimes the same structure is used as a delay-chain, because the sample `x` takes 4 updates to travel from `shr[0]` to `shr[3]`. This is a very common element in hardware designs.

Fig. 2.5 shows the RTL for this design, as expected the `for` has been unrolled, thus all the summing is done.




Fig. 2.5: Synthesis result of Listing 2.5 (Intel Quartus RTL viewer)

Optimizing the design

This design can be made generic by changing the `__init__` function to take the window length as a parameter (Listing 2.6).

Listing 2.6: Generic sliding adder

```
class SlidingAdder(HW):
    def __init__(self, window_len):
        self.shr = [0] * window_len
    ...
```

The problem with this design is that it starts using more resources as the `window_len` gets larger as every stage requires a separate adder. Another problem is that the critical path gets longer, decreasing the clock rate. For example, the design with `window_len=4` synthesises to maximum clock of 170 MHz, while `window_len=6` to only 120 MHz.

Todo

MHz on what FPGA?




Fig. 2.6: RTL of `window_len=6`, the red line shows the critical path (Intel Quartus RTL viewer)

In that sense, it can be considered a poor design, as it is hard to reuse. Conveniently, the algorithm can be optimized to use only 2 adders, no matter the window length. Listing 2.7 shows that instead of summing all the elements, the overlapping part of the previous calculation can be used to significantly optimize the algorithm.

Listing 2.7: Optimizing the sliding adder algorithm by using recursive implementation

```
y[4] = x[4] + x[5] + x[6] + x[7] + x[8] + x[9]
y[5] =      x[5] + x[6] + x[7] + x[8] + x[9] + x[10]
y[6] =      x[6] + x[7] + x[8] + x[9] + x[10] + x[11]

# reusing overlapping parts implementation
y[5] = y[4] + x[10] - x[4]
y[6] = y[5] + x[11] - x[5]
```

Listing 2.8 gives the implementation of the optimal sliding adder; it features a new reg-

ister sum`, that keeps track of the previous output. Note that the ``shr stayed the same, but is now rather used as a delay-chain.

Listing 2.8: Optimal sliding adder

```
class OptimalSlideAdd(HW):
    def __init__(self, window_len):
        self.shr = [0] * window_len
        self.sum = 0

        self._delay = 1

    def main(self, x):
        self.next.shr = [x] + self.shr[:-1]

        self.next.sum = self.sum + x - self.shr[-1]
        return self.sum
    ...
```

Fig. 2.7 shows the synthesis result; as expected, the critical path is along 2 adders.




Fig. 2.7: Synthesis result of Listing 2.5, window_len=4 (Intel Quartus RTL viewer)

Simulations results (Fig. 2.8) show that the hardware desing behaves exactly as the software model. Note that the class has `self._delay=1` to compensate for the register delay.




Fig. 2.8: Simulation results for OptimalSlideAdd(window_len=4)

2.3 Fixed-point designs

Examples in the previous chapters have used only the `integer` type, in order to simplify the designs.

Todo

explain why float costs greatly?

DSP algorithms are mostly described using floating point numbers. As shown in previous sections, every operation in hardware takes resources and floating point calculations cost greatly. For that reason, fixed-point arithmetic is often used in hardware designs.

Fixed-point arithmetic is in nature equal to integer arithmetic and thus can use the DSP blocks that come with many FPGAs (some high-end FPGAs have also floating point DSP blocks [14]).

2.3.1 Basics

Pyha defines `Sfix` for FP objects; it is a signed number. It works by defining bits designated for `left` and `right` of the decimal point. For example `Sfix(0.3424, left=0, right=-17)` has 0 bits for integer part and 17 bits for the fractional part. Listing 2.9 shows some examples. more information about the fixed point type is given on APPENDIX.

Todo

Add more information about fixed point stuff to the appendix

Listing 2.9: Example of `Sfix` type, more bits give better results

```
>>> Sfix(0.3424, left=0, right=-17)
0.34239959716796875 [0:-17]
>>> Sfix(0.3424, left=0, right=-7)
0.34375 [0:-7]
>>> Sfix(0.3424, left=0, right=-4)
0.3125 [0:-4]
```

The default FP type in Pyha is `Sfix(left=0, right=-17)`, it represents numbers between `[-1;1]` with resolution of `0.000007629 (2**-17)`. This format is chosen because it fits into common FPGA DPS blocks (18 bit signals [15]) and it can represent normalized numbers.

The general recommendation is to keep all the inputs and outputs of the block in the default type.

2.3.2 Fixed-point sliding adder

Consider converting the sliding window adder, described in [Section 2.2.2](#), to FP implementation. This requires changes only in the `__init__` function ([Listing 2.10](#)).

Listing 2.10: Fixed-point sliding adder

```
def __init__(self, window_size):
    self.shr = [Sfix()] * window_size
    self.sum = Sfix(left=0)
    ...
```

The first line sets `self.shr` to store `Sfix()` elements. Notice that it does not define the fixed-point bounds, meaning it will store ‘whatever’ is assigned to it. The final bounds are determined during simulation.

Todo

lazy stuff needs more explanation

The `self.sum` register uses another lazy statement of `Sfix(left=0)`, meaning that the integer bits are forced to 0 bits on every assign to this register. The fractional part is left determined by simulation. The rest of the code is identical to the one described in [Section 2.2.2](#).

Synthesis results are shown in [Fig. 2.9](#). In general, the RTL diagram looks similar to the one at [Section 2.2.2](#). First noticeable change is that the signals are now 18 bits wide due to the default FP type. The second addition is the saturation logic, which prevents the wraparound behaviour by forcing the maximum or negative value when they are out of fixed point format. Saturation logic is by default enabled for FP types.




Fig. 2.9: RTL of fixed-point sliding adder (Intel Quartus RTL viewer)

[Fig. 2.10](#) plots the simulation results for input of random signal in $[-0.5;0.5]$ range. Notice that the hardware simulations are bounded to $[-1;1]$ range by the saturation logic, that is why the model simulation is different at some points.




Fig. 2.10: Simulation results of FP sliding sum

Simulation functions can automatically convert ‘floating-point’ inputs to default FP type. In same manner, FP outputs are converted to floating point numbers. That way, the designer does not have to deal with FP numbers in unit-testing code. An example is given in [Listing 2.11](#).

Listing 2.11: Test fixed-point design with floating-point numbers

```
dut = OptimalSlidingAddFix(window_len=4)
x = np.random.uniform(-0.5, 0.5, 64)
y = simulate(dut, x)
# plotting code ...
```

2.4 Summary

This chapter has demonstrated that in Pyha traditional software language features can be used to infer hardware components and their outputs are equivalent. One must still keep in mind how the code converts to hardware, for example that the loops will be unrolled. A major difference between hardware and software is that in hardware, every arithmetical operator takes up resources.

Class variables can be used to add memory to the design. In Pyha, class variables must be assigned to `self.next` as this mimics the **delayed** nature of registers. The general rule is to always register the outputs of Pyha designs.

DSP systems can be implemented by using the fixed-point type. Pyha has ‘semi-automatic conversion’ from floating point to fixed point numbers. Verifying against floating point model helps the design process.

Reusing Pyha designs is easy thanks to the object-oriented style that also works well for design abstraction.

Pyha provides the `simulate` function that can automatically run Model, Pyha, RTL and GATE level simulations. In addition, `assert_simulate` can be used for fast design of unit-

tests. These functions can automatically handle fixed point conversion, so that tests do not have to include fixed point semantics. Pyha designs are debuggable in the Python domain.

In Pyha all class variables are interpreted as hardware registers. The `__init__` function may contain any Python code to evaluate reset values for registers.

The key difference between software and hardware approaches is that hardware registers have **delayed assignment**, they must be assigned to `self.next`.

The delay introduced by the registers may drastically change the algorithm, that is why it is important to always have a model and unit tests, before starting hardware implementation. The model delay can be specified by `self._delay` attribute, this helps the simulation functions to compensate for the delay.

Registers are also used to shorten the critical path of chained logic elements, thus allowing higher clock rate. It is encouraged to register all the outputs of Pyha designs.

In Pyha, DSP systems can be implemented by using the fixed-point type. The combination of ‘lazy’ bounds and default Sfix type provide simplified conversion from floating point to fixed point. In that sense it could be called ‘semi-automatic conversion’.

Simulation functions can automatically perform the floating to fixed point conversion, this enables writing unit-tests using floating point numbers.

Comparing the FP implementation to the floating-point model can greatly simplify the final design process.

Chapter 3

Conversion to VHDL

This chapter shows how Pyha converts to VHDL.

Todo

Pilt converterist

- Simulatsioon tüüptide leidmiseks
 - Python syntax to VHDL
 - Sequential OOP VHDL IR
-

First part of this chapter introduces the Sequential OOP VHDL IR.

While other high level tools convert to very low-level VHDL, then Pyha takes a different approach by first developing a feasible model in VHDL and then using Python to get around VHDL ugly parts.

Many tools on the market are capable of converting higher level language to VHDL. However, these tools only make use of the very basic dataflow semantics of VHDL language, resulting in a complex conversion process and typically unreadable VHDL output.

The design choices done in the process of Pyha design have focused on simplicity. The conversion process of Python code to VHDL is straight-forward as the synthesis tools are already capable of elaborating sequential VHDL code. This work contributes the object-oriented VHDL design way that allows defining registers in sequential code. Thanks to that, the OOP Python code can be simply mapped to OOP VHDL code. Result is readable (keeps hierarchy) VHDL code that may provide a bridge for people that already know VHDL.

3.1 Sequential, Object-oriented style for VHDL

This chapter develops sequential synthesizable object-oriented (OOP) programming model for VHDL. The main motivation is to use it as an intermediate language for High-Level synthesis of hardware.

VHDL has been chosen over SystemVerilog(SV) because it is a strict language and forbids many mistakes during compile time. SV on the other hand is much more permissive, for example allowing out-of-bounds array indexing [16].

Sequential programming in VHDL has been studied by Jiri Gaisler in [17]. He showed that combinatory logic is easily described by fully sequential functions. He proposed the ‘two-process’ design method, where one of the processes is for comb and other for registers. His work is limited to one clock domain.

This section’s contribution is the extension of the ‘two process’ model by adding an Object-oriented approach. The basic idea of OOP is to bundle data and define functions that perform actions on this data. This idea fits well with hardware design, as ‘data’ can be thought as registers and combinatory logic as functions that perform operations on the data.

VHDL has no direct support, but the OOP style can be still used by combining data in records (same as ‘C’ struct) and passing them as parameters to functions. This is essentially the same way how C programmers do it.

Listing 3.1 demonstrates pipelined multiply-accumulate(MAC), written in OOP VHDL. Recall that all the items in the `self_t` are to be registers. One inconvenience is that VHDL procedures cannot ‘return’, instead ‘out’ direction arguments must be used. On the other hand this helps to handle Python functions that can return multiple values.

Listing 3.1: OOP style multiply-accumulate in VHDL

```
type self_t is record
    mul: integer;
    acc: integer;
    coef: integer;
end record;

procedure main(self: inout self_t; a: in integer; ret_0: out integer) is
begin
    self.mul := a * self.coef;
    self.acc := self.acc + self.mul;
    ret_0 := self.acc;
end procedure;
```

The synthesis results (Fig. 3.1) show that a functionally correct MAC has been implemented. However, in terms of hardware, it is not quite what was wanted. The data model specified 3 registers, but only the one for ‘acc’ is present and even this is at the wrong location.

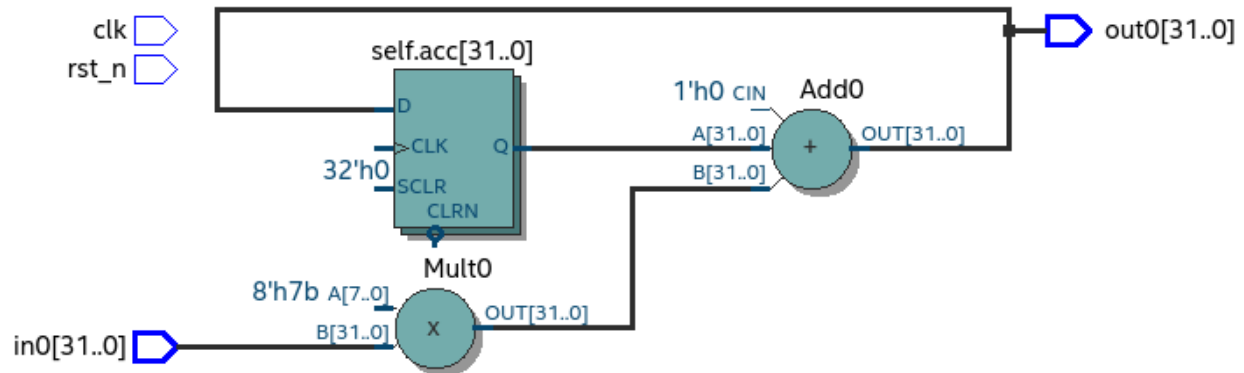


Fig. 3.1: Unexpected synthesis result of Listing 3.1 (Intel Quartus RTL viewer)

3.1.1 Defining registers with variables

Clearly the way of defining registers is not working properly. The mistake was to expect that the registers work in the same way as ‘class variables’ in traditional programming languages.

Hardware registers have just one difference to class variables, the value assigned to them does not take effect immediately, but rather on the next clock edge. That is the basic idea of registers, they take a new value on clock edge. When the value is set at **this** clock edge, it will be taken on **next** clock edge.

Trying to stay in the software world, we can abstract away the clock edge by thinking that it denotes the call to the ‘main’ function. Meaning that registers take the assigned value on the next function call, meaning assignment is delayed by one function call.

VHDL defines a special assignment operator for this kind of delayed assignment, it is called ‘signal assignment’. It must be used on VHDL signal objects like `a <= b`.

VHDL signals really come down to just having two variables, to represent the **next** and **current** values. Signal assignment operator sets the value of **next** variable. On the next simulation delta, **current** is automatically set to equal **next**.

This two variable method has been used before, for example Pong P. Chu, author of one of the most reputed VHDL books, suggests to use this style in defining sequential logic in VHDL [18]. The same semantics are also used in MyHDL signal objects [19].

Adapting this style for the OOP data model is shown in Listing 3.2. The new data model extends the structure to include the ‘nexts’ object, that can used to assign **next** value for registers, for example `self.nexts.acc := 0`.

Listing 3.2: Data model with **next**, in OOP-style VHDL

```
type next_t is record
  mul: integer;
  acc: integer;
```

```

    coef: integer;
end record;

type self_t is record
    mul: integer;
    acc: integer;
    coef: integer;

    nexts: next_t; -- new element to hold 'next state' value
end record;

procedure main(self: inout self_t; a: integer; ret_0: out integer) is
begin
    self.nexts.mul := a * self.coef;          -- now assigns to self.nexts
    self.nexts.acc := self.acc + self.mul;    -- now assigns to self.nexts
    ret_0 := self.acc;
end procedure;

```

Now the loading of **next** to **current** must now be done manually. Listing 3.3 defines new function ‘update_registers’, taking care of this task.

Listing 3.3: Function to update registers, in OOP-style VHDL

```

procedure update_register(self: inout self_t) is
begin
    self.mul := self.nexts.mul;
    self.acc := self.nexts.acc;
    self.coef:= self.nexts.coef;
end procedure;

```

Note: Function ‘update_registers’ is called on clock raising edge. While the ‘main’ is called as combinatory function.

Todo

add simple top level example here?

Synthesising this results in expected logic, that is MAC with pipelined registers (Fig. 3.2).




Fig. 3.2: Synthesis result of the revised code (Intel Quartus RTL viewer)

3.1.2 Creating instances

Todo

consider removing this section, quite useless..

The general approach of creating instances is to define new variables of the ‘self.t’ type, Listing 3.4 gives an example of this.

Listing 3.4: Class instances by defining records, in OOP-style VHDL

```
variable mac0: MAC.self_t;
variable mac1: MAC.self_t;
```

The next step is to initialize the variables, this can be done at the variable definition, for example: `variable mac0: self_t := (mul=>0, acc=>0, coef=>123, nexts=>(mul=>0, acc=>0, coef=>123));`

The problem with this method is that all data-model must be initialized (including ‘nexts’), this will get unmaintainable very quickly, imagine having an instance that contains another instance or even array of instances. In some cases it may also be required to run some calculations in order to determine the initial values.

Traditional programming languages solve this problem by defining class constructor, executing automatically for new objects.

In the sense of hardware, this operation can be called ‘reset’ function. Listing 3.5 is a reset function for the MAC circuit. It sets the initial values for the data model and can also be used when reset signal is asserted.

Listing 3.5: Reset function for MAC, in OOP-style VHDL

```
procedure reset(self: inout self_t) is
begin
    self.nexts.coef := 123;
    self.nexts.mul  := 0;
    self.nexts.sum  := 0;
```

```
    update_registers(self);  
end procedure;
```

But now the problem is that we need to create a new reset function for each instance.

This can be solved by using VHDL ‘generic packages’ and ‘package instantiation declaration’ semantics [20]. Package in VHDL just groups common declarations to one namespace.

In case of the MAC class, the ‘coef’ reset value could be set as package generic. Then each new package initialization could define new reset value for it (Listing 3.6).

Listing 3.6: Initialize new package MAC_0, with ‘coef’ 123

```
package MAC_0 is new MAC  
    generic map (COEF => 123);
```

Unfortunately, these advanced language features are not supported by most of the synthesis tools. A workaround is to either use explicit record initialization (as at the start of this chapter) or manually make new package for each instance.

Both of these solutions require unnecessary workload.

The Python to VHDL converter (developed in the next chapter), uses the later option, it is not a problem as everything is automated.

3.1.3 Final OOP model

Currently the OOP model consists of following elements:

- Record for ‘next’
- Record for ‘self’
- User defined functions (like ‘main’)
- ‘Update registers’ function
- ‘Reset’ function

VHDL supports ‘packages’ to group common types and functions into one namespace.

Listing 3.7 shows the template package for VHDL ‘class’. All the class functionality is now in common namespace.

Listing 3.7: Package template for OOP style VHDL

```
package Class is  
    type next_t is record  
        ...  
    end record;
```

```

    type self_t is record
        ...
        nexts: next_t;
    end record;

    -- function prototypes
end package;

package body Class is
    procedure reset(self: inout self_t) is
        ...
    procedure update_registers(self: inout self_t) is
        ...
    procedure main(self: inout self_t) is
        ...
        -- other user defined functions
    end package body;

```

3.1.4 Examples

Creating a new class that connects two MAC instances in series is simple, first we need to create two MAC packages called MAC_0 and MAC_1 and add them to the data model (Listing 3.8). The next step is to call MAC_0 operation on the input and then pass the output through MAC_1, whose output is the final output.

Todo

why MAC_0 and MAC_1?

Listing 3.8: Series MACs in OOP-style VHDL

```

type self_t is record
    mac0: MAC_0.self_t; -- define 2 MACs as part of data model
    mac1: MAC_1.self_t;

    nexts: next_t;
end record;

procedure main(self: inout self_t; a: integer; ret_0: out integer) is
    variable out_tmp: integer;
begin
    MAC_0.main(self.mac0, a, ret_0=>out_tmp);           -- connect MAC_0 output to_
    ↪MAC_1 input

```

```

MAC_1.main(self.mac1, out_tmp, ret_0=>ret_0);
end procedure;

```

Synthesis result shows that two MACs are connected in series Fig. 3.3.




Fig. 3.3: Synthesis result of the new class (Intel Quartus RTL viewer)

Connecting two MAC's instead in parallel can be done with simple modification to 'main' function, that instead now returns both outputs (Listing 3.9).

Listing 3.9: Main function for parallel instances, in OOP-style VHDL

```

procedure main(self: inout self_t; a: integer; ret_0: out integer; ret_1: out
integer) is
begin
    MAC_0.main(self.mac0, a, ret_0=>ret_0); -- return MAC_0 output
    MAC_1.main(self.mac1, a, ret_0=>ret_1); -- return MAC_1 output
end procedure;

```

Two MAC's are synthesized in parallel, as shown in Fig. 3.4.

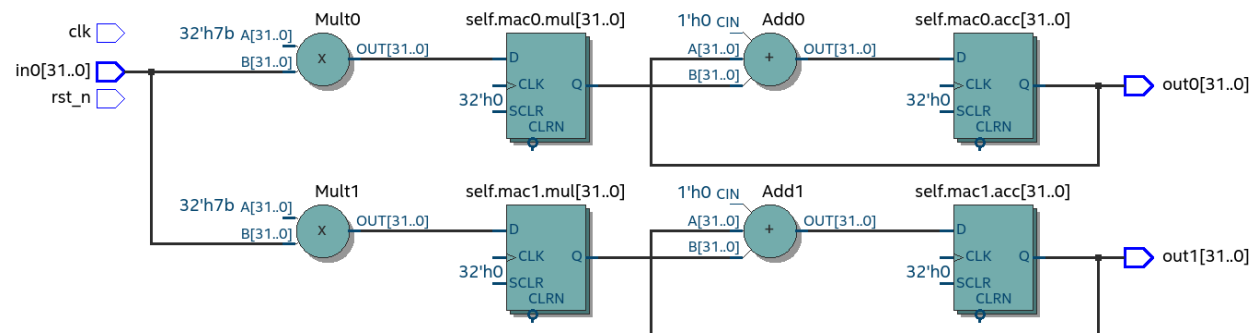


Fig. 3.4: Synthesis result of Listing 3.9 (Intel Quartus RTL viewer)

3.1.5 Multiple clock domains

Multiple clock domains can be easily supported by updating registers at different clock domains. By reusing the parallel MAC's example, consider that MAC_0 and MAC_1 work in different clock domain. For this only the top level process must be modified (Listing 3.10), rest of the code stays the same.

Listing 3.10: Top-level for multiple clocks, in OOP-style VHDL

```

if (not rst_n) then
    ReuseParallel_0.reset(self);
else
    if rising_edge(clk0) then
        MAC_0.update_registers(self.mac0); -- update 'mac0' on 'clk0' rising edge
    end if;

    if rising_edge(clk1) then
        MAC_1.update_registers(self.mac1); -- update 'mac1' on 'clk1' rising edge
    end if;
end if;

```

Synthesis result (Fig. 3.5) show that registers are clocked by different clocks. The reset signal is common for the whole design.




Fig. 3.5: Synthesis result with modified top-level process (Intel Quartus RTL viewer)

3.2 Converting Python to VHDL

The Python to VHDL conversion process relies heavily on the results of last chapter, that allows sequential OOP Python code easily map to VHDL. Even so, converting Python syntax to VHDL poses some problems.

The biggest challenge in conversion from Python to VHDL is types, namely Python does not have them, while VHDL has. Conversion process must find all the types for Python variables, the process of this is described in.

After the types are all known, the design can be converted from Python to VHDL syntax. This requires some way of traversing the Python source code and applying VHDL rated transforms.

Conversion progress requires no understanding of the source code nor big modifications. ...
_pyvhdl_types:

3.2.1 Finding the types

Python is dynamically typed language, meaning that types come into play only when the code is running. On the other hand VHDL is statically typed, all the types must be written in source code.

The advantage of the Python way is that it is easier to program, no need to define variables and ponder about the types. Downsides are that there may be unexpected bugs when some variable changes type. In some cases dynamic typing may also reduce code readability.

In sense of conversion, dynamic typing poses a major problem, somehow the missing type info should be recovered for the VHDL code. Most straightforward way to solve this is to try finding the variables value from code, for example `a = 5`, clearly type of `a` is integer. Problem with this method is that is much more complex than it initially appears. For example `a = b`. To find the type of `a` converter would need to lookup type of `b`, these kind of suffixes can get really complex.

Alternative, and what Pyha is using, is to run the Python code so all the variables get some value, the value can be inspected programmically and type inferred. For example, consider the class on [Listing 3.11](#).

Listing 3.11: Example Python class, what are the types?

```
class SimpleClass(HW):
    def __init__(self, coef):
        self.coef = coef

    def main(self, a):
        local_var = a
```

[Listing 3.12](#) show example for getting the type of class variable. It initializes the class with argument 5, that is assigned to the `coef` variable. Then `type()` can be used to query the variable type. On the example result is `int`, so this can be converted to VHDL `integer` type.

Listing 3.12: Using `type()` to get type name

```
>>> dut = SimpleClass(5)
>>> dut.coef
5
>>> type(dut.coef)
<class 'int'>
```

Pyha deduces registers initial values in same way, only the first assigned value is considered.

Local variables, like `local_var` and argument `a` on [Listing 3.12](#) are harder to deduce as Python provides no way of accessing function locals scope. Note that locals exist only in the stack, thus after the function call they are lost forever. Luckily this problem has been

encountered before in [21], which ‘hacks’ the Python profiling interface in order to save the locals for each function. Pyha uses this approach to keep track of the local values.

Listing 3.13: Function locals variable type

```
>>> dut.main.locals # before any call, locals are unknown
{}
>>> dut.main(1) # call function
>>> dut.main.locals # locals can be extracted
{'a': 1, 'local_var': 1}
>>> type(dut.main.locals['local_var'])
<class 'int'>
```

Advantage of this method is low complexity, another perk is that this way could be used to keep track of all the variable values, in future this can enable the automatic conversion from floating point to fixed point. In addition, this way allows the ‘lazy’ coding, for example where fixed-point gains the bound limit only during the execution of the design.

Downside is that each function in the design must be executed before conversion is possible. Also the conversion result may depend on the data types that are inputted to the functions, but this can also be an advantage.

3.2.2 Syntax conversion

The syntax of Python and VHDL is surprisingly similar. VHDL is just much more verbose, requires types and Python has indention oriented blocks.

Python provides some tools that simplify the traversing of source files, like abstract syntax tree (AST) module and lib2to3. These tools work by parsing the Python file into a tree structure, that can be then traversed and modified. For example the MyHDL conversion is based on this. This method works but is quite complex and requires alot of code.

Lately new project has emerged called RedBaron [22], that aims to simplify operations with Python source code. It features rich tools for searching and modifying the source code. Unlike AST it also keeps all the formatting in the code, including comments. RedBaron parses the source code into rich objects, for example the `a = 5` would result in a `AssignmentNode` object that has an `__str__` function that instruct how these kind of objects are written out.

Pyha overwrites the `__str__` method to instead of `= print :=` and also add `;` to the end of statement. Resulting in a VHDL compatible statement `a := 5;`. Beauty of this is that this simple modification actually turns **all** the Python style assignments to VHDL style.

Listing 3.14 shows a more complex Python code that is converted to VHDL (Listing 3.15), by Pyha. Most of the transforms are obtained by the same method described above. Some of the transforms are a bit more complex, like figuring out what variables need to be defined in VHDL code.

Listing 3.14: Python function to be converted to VHDL

```
def main(self, x):  
    y = x  
    for i in range(4):  
        y = y + i  
  
    return y
```

Listing 3.15: Conversion of Listing 3.14 assuming integer types

```
procedure main(self:inout self_t; x: integer; ret_0:out integer) is  
    variable y: integer;  
begin  
    y := x;  
    for i in 0 to (4) - 1 loop  
        y := y + i;  
    end loop;  
  
    ret_0 := y;  
end procedure;
```

3.3 Summary

The sequential object-oriented VHDL model is one of the contributions of this thesis. It has been developed to provide simpler conversion from Python to VHDL. Pyha converts directly to the VHDL model by using RedBaron based syntax conversions. Type information is required through the simulation before conversion.

Like HLS must do much work to deduce registers.. Pyha can convert basically line by line, very simple.

Chapter 4

Case studies

4.1 Moving average filter

Todo

rephrase as this is copy paste!

The moving average (MA) is the most common filter in DSP, mainly because it is the easiest digital filter to understand and use. In spite of its simplicity, the moving average filter is optimal for a common task: reducing random noise while retaining a sharp step response. This makes it the premier filter for time domain encoded signals [23].

Fig. 4.1 shows that MA is a good algorithm for noise reduction. Increasing the window length reduces more noise but also increases the complexity and delay of the system (MA is a special case of FIR filter, same delay semantics apply).




Fig. 4.1: MA algorithm in removing noise

Good noise reduction performance can be explained by the frequency response of MA (Fig.

4.2), showing that it is a low-pass filter. Passband width and stopband attenuation are controlled by the window length.




Fig. 4.2: Frequency response of MA filter

MA is implemented by using a sliding sum that is divided by the sliding window length. The sliding sum part has already been implemented in [Section 2.3.2](#). The division can be implemented by a shift right operation, assuming that the divisor is power of two.

In addition, division can be performed on each sample instead of on the sum, that is $(a + b) / c = a/c + b/c$. Doing this guarantees that the `sum` variable is always in the $[-1;1]$ range, thus the saturation logic can be removed.

Listing 4.1: MA implementation in Pyha

```

1 class MovingAverage(HW):
2     def __init__(self, window_len):
3         self.window_pow = Const(int(np.log2(window_len)))
4
5         self.shr = [Sfix()] * window_len
6         self.sum = Sfix(0, 0, -17, overflow_style=fixed_wrap)
7         self._delay = 1
8
9     def main(self, x):
10        div = x >> self.window_pow
11
12        self.next.shr = [div] + self.shr[:-1]
13        self.next.sum = self.sum + div - self.shr[-1]
14        return self.sum
15    ...

```

The code in [Listing 4.1](#) makes only few changes to the sliding sum:

- On line 3, `self.window_pow` stores the bit shift count (to support generic `window_len`)
- On line 6, type of `sum` is changed so that saturation is turned off

- On line 10, shift operator performs the division

Fig. 4.3 shows the synthesized result of this work; as expected it looks very similar to the sliding sum RTL schematics. In general, shift operators are hard to notice on the RTL schematics because they are implemented by routing semantics.




Fig. 4.3: RTL view of moving average (Intel Quartus RTL viewer)

MA is an optimal solution for performing matched filtering of rectangular pulses [23]. This is important for communication systems. Fig. 4.4 shows an example of a digital signal, that is corrupted with noise. MA with window length equal to samples per symbol can recover the signal from the noise.




Fig. 4.4: Moving average as matched filter

The ‘model’ deviates from rest of the simulations because the input signal violates the $[-1;1]$ bounds and hardware simulations are forced to saturate the values.

4.2 Linear-phase DC removal Filter

Pyha has been designed in the way that it can represent RTL designs exactly as the user defines, however thanks to the object-oriented nature all these low level details can be abstracted away and then Pyha turns into HLS language. To increase productivity, abstraction is needed.

Pyha is based on the object-oriented design practices, this greatly simplifies the design reuse as the classes can be used to initiate objects. Another benefit is that classes can abstract away the implementation details, in that sense Pyha can become a high-level synthesis (HLS) language.

This chapter gives an example on how to reuse the moving average filter for ...

Direct conversion (homodyne or zero-IF) receivers have become very popular recently especially in the realm of software defined radio. There are many benefits to direct conversion receivers, but there are also some serious drawbacks, the largest being DC offset and IQ imbalances [24].

DC offset looks like a peak near the 0Hz on the frequency response. In the time domain, it manifests as a constant component on the harmonic signal.

In [25], Rick Lyons investigates the use of moving average algorithm as a DC removal circuit. This works by subtracting the MA output from the input signal. The problem of this approach is the 3 dB passband ripple. However, by connecting multiple stages of MA's in series, the ripple can be avoided (Fig. 4.5) [25].




Fig. 4.5: Frequency response of DC removal filter (MA window length is 8)

The algorithm is composed of two parts. First, four MA's are connected in series, outputting the DC component of the signal. Second, the MA's output is subtracted from the input signal, thus giving the signal without DC component. Listing 4.2 shows the Pyha implementation.

Listing 4.2: DC-Removal implementation

```
class DCRemoval(HW):
    def __init__(self, window_len):
        self.mavg = [MovingAverage(window_len), MovingAverage(window_len),
                     MovingAverage(window_len), MovingAverage(window_len)]
        self.y = Sfix(0, 0, -17)

        self._delay = 1

    def main(self, x):
        # run input signal over all the MA's
        tmp = x
        for mav in self.mavg:
            tmp = mav.main(tmp)

        # dc-free signal
        self.next.y = x - tmp
        return self.y

    ...
```

This implementation is not exactly following that of [25]. They suggest to delay-match the step 1 and 2 of the algorithm, but since we can assume the DC component to be more or less stable, this can be omitted.

Fig. 4.6 shows that the synthesis generated 4 MA filters that are connected in series, output of this is subtracted from the input.




Fig. 4.6: Synthesis result of DCRemoval(window_len=4) (Intel Quartus RTL viewer)

In a real application, one would want to use this component with larger `window_len`. Here 4 was chosen to keep the RTL simple. For example, using `window_len=64` gives much better cutoff frequency (Fig. 4.7); FIR filter with the same performance would require hundreds of taps [25]. Another benefit is that this filter delays the signal by only 1 sample.




Fig. 4.7: Comparison of frequency response

This implementation is also very light on the FPGA resource usage ([Listing 4.3](#)).

Listing 4.3: Cyclone IV FPGA resource usage for DCRemoval(window_len=64)

| | |
|------------------------------------|-----------------------------|
| Total logic elements | 242 / 39,600 (< 1 %) |
| Total memory bits | 2,964 / 1,161,216 (< 1 %) |
| Embedded Multiplier 9-bit elements | 0 / 232 (0 %) |

[Fig. 4.8](#) shows the situation where the input signal is corrupted with a DC component (+0.25), the output of the filter starts countering the DC component until it is removed.




Fig. 4.8: Simulation of DC-removal filter in the time domain

4.3 Comparison to other tools

MyHDL is following the event-driven approach which is a trait of the classical HDL's. It features an function based design that is very similar to Verilog processes. In general the synthesizable subset of MyHDL is very limited, it has been found that the tool is more useful for high-level modeling purposes [\[26\]](#). Another package in the Python ecosystem is Migen, that replaces the event-driven paradigm with the notions of combinatorial and synchronous statements [\[27\]](#). Migen can be considered as meta-programming in Python so it is a bit complicated. Both Migen and MyHDL are more aimed at the control logic, neither implements the fixed-point data type, that is a standard for hardware DSP designs.

Pyha aims to raise the abstraction level by using sequential object-oriented style, major advantage of this is that existing blocks can be connected together in purely Pythonic way, the designer needs to know nothing about the underlying RTL implementation.

Chapter 5

Summary

This work introduced new Python based HDL language called Pyha. That is an sequential object-oriented HDL language. Overview of Pyha features have been given and shown that the tool is suitable to use with model based DSP systems. It was also demonstrated that Pyha supports fixed point types and semi-automatic onversion from floating point types. Pyha also provides good support for unit test and designs are debuggable.

Two case studies were presented, first the moving average filter that was implemented in RTL level. Second example demonstrated that blocks written in Pyha can be reused in pure Python way, developed linear phase DC removal filter by reusing the implementation of moving average filter. Synthesisability has been demonstrated on Cyclone IV device (BladeRF).

Another contribution of this thesis is the sequential object-oriented VHDL model. This was developed to enable simple conversion of Pyha to VHDL. One of the advantages of this work compared to other tools is the simplicity of how it works.

Future perspectives are to implement more DSP blocks, especially by using GNURadio blocks as models. That may enable developing an work-flow where GNURadio designs can easily be converted to FPGA. In addition, the Pyha system could be improved to add automatic fixed point conversion and support for multiple clock domains.

In this work Pyha has been designed to add DSP related features to the Python conversion scope, this includes fixed-point type and semi-automatic conversion from floating point. In addition Pyha integrates the model to designs and test functions simplification. Pyha includes functions that help verification by automatically running all the simulations, asserting that model is equivalent to the synthesis result, tests defined for model can be reused for RTL, model based verifaciton. Pyha designs are also simulatable and debuggable in Python domain. The design of Pyha also supports fully automatic conversion but currently this is left as a future work.

Pyha is a fully sequential language that works on purely Python code. However Pyha resides in the RTL level, allowing to define each and every register. In that sende Pyha is at somewhere between the HLS and HDL language. Pyha aims to raise the abstraction level

by using the object-oriented style, so that the RTL details can be easily abstracted away. One major advantage of Pyha is that existing blocks can be connected together in purely Python way, the designer needs to know nothing about hardware design or underlying RTL implementation.

The design choices done in the process of Pyha design have focused on simplicity. The conversion process of Python code to VHDL is straight-forward as the synthesis tools are already capable of elaborating sequential VHDL code. This work contributes the object-oriented VHDL desing way that allows defining registers in sequential code. Thanks to that, the OOP Python code can be simply mapped to OOP VHDL code. Result is readable (keeps hirarchy) VHDL code that may provide an bridge for people that already know VHDL.

Limited to one clock domain? In some seneses the Python part could be considers as Python binding to VHDL OOP model. Convert to HLS langauge instead of VHDL, then the designer could choose to to either design for RTL or HLS, this is more as an futures perspective, this thesis works only with the RTL part.

Long term goal of the project is to develop enough blocks that are functionally equal to GNURadio blocks, so that flow-graphs could be converted to FPGA designs, thus providing an open-source alternative for Simulink based flows.

Bibliography

- [1] Richard Goering. Designers debate ‘vhdl is dead’ assertion. URL: http://www.eetimes.com/document.asp?doc_id=1216820.
- [2] Ieee p1076 working group vhdl analysis and standardization group (vasg). URL: <http://www.eda-twiki.org/cgi-bin/view.cgi/P1076/WebHome>.
- [3] Open source vhdl verification methodology (osvvm). URL: <http://osvvm.org/>.
- [4] Lars Asplund. Vunit. URL: <http://vunit.github.io/>.
- [5] Myhdl. URL: <http://www.myhdl.org>.
- [6] Clash. URL: <http://www.clash-lang.org/>.
- [7] Jonathan Bachrach. Chisel: constructing hardware in a scala embedded language. 2012.
- [8] Wim Meeus, Kristof Van Beeck, Toon Goedemé, Jan Meel, and Dirk Stroobandt. An overview of today’s high-level synthesis tools. *Design Automation for Embedded Systems*, 16(3):31–51, 2012. URL: <http://dx.doi.org/10.1007/s10617-012-9096-8>, doi:10.1007/s10617-012-9096-8.
- [9] F. Winterstein, S. Bayliss, and G. A. Constantinides. High-level synthesis of dynamic data structures: a case study using vivado hls. In *2013 International Conference on Field-Programmable Technology (FPT)*, 362–365. Dec 2013. doi:10.1109/FPT.2013.6718388.
- [10] Hayden Kwok-Hay So. *BORPH: An Operating System for FPGA-Based Reconfigurable Computers*. PhD thesis, UNIVERSITY OF CALIFORNIA, BERKELEY, 2007.
- [11] Robert Ghilduta and Brian Padalino. Bladerf vhdl ads-b decoder. URL: <https://www.nuand.com/blog/bladerf-vhdl-ads-b-decoder/>.
- [12] Neil Lawrence. Gpy: moving from matlab to python. URL: <http://inverseprobability.com/2013/11/25/gpy-moving-from-matlab-to-python>.
- [13] Pierre Carboneille. Pypl popularity of programming language. URL: <http://pypl.github.io/PYPL.html>.
- [14] Amulya Vishwanath. Enabling high-performance floating-point designs. URL: https://www.altera.com/content/dam/altera-www/global/en_US/pdfs/literature/wp/wp-01267-fpgas-enable-high-performance-floating-point.pdf.

-
- [15] Altera. Cyclone iv fpga device family overview. 2016.
 - [16] Aurelian Ionel Munteanu. Gotcha: access an out of bounds index for a systemverilog fixed size array. URL: <http://www.amiq.com/consulting/2016/01/26/gotcha-access-an-out-of-bounds-index-for-a-systemverilog-fixed-size-array/>.
 - [17] Jiri Gaisler. A structured vhdl design method. URL: <http://www.gaisler.com/doc/vhdl2proc.pdf>.
 - [18] Pong P. Chu. *RTL Hardware Design Using VHDL: Coding for Efficiency, Portability, and Scalability*. Wiley, 2006.
 - [19] Jan Decaluwe. Why do we need signal assignments? URL: <http://www.jandecaluwe.com/hdl2proc/signal-assignments.html>.
 - [20] IEEE. Ieee standard vhdl language reference manual. *IEEE Std 1076-2008 (Revision of IEEE Std 1076-2002)*, 2009.
 - [21] Pietro Berkes and Andrea Maffezoli. Decorator to expose local variables of a function after execution. URL: <http://code.activestate.com/recipes/577283-decorator-to-expose-local-variables-of-a-function-/>.
 - [22] Laurent Peuch. Redbaron: bottom-up approach to refactoring in python. URL: <http://redbaron.pycqa.org/>.
 - [23] Steven W. Smith. *The Scientist & Engineer's Guide to Digital Signal Processing*. California Technical Pub, 1997.
 - [24] BladeRF community. Dc offset and iq imbalance correction. 2017. URL: <https://github.com/Nuand/bladeRF/wiki/DC-offset-and-IQ-Imbalance-Correction>.
 - [25] Rick Lyons. Linear-phase dc removal filter. 2008. URL: <https://www.dsprelated.com/showarticle/58.php>.
 - [26] Jan Decaluwe. It's a simulation language! URL: <http://www.jandecaluwe.com/blog/its-a-simulation-language.html>.
 - [27] Migen. URL: <https://m-labs.hk/gateway.html>.